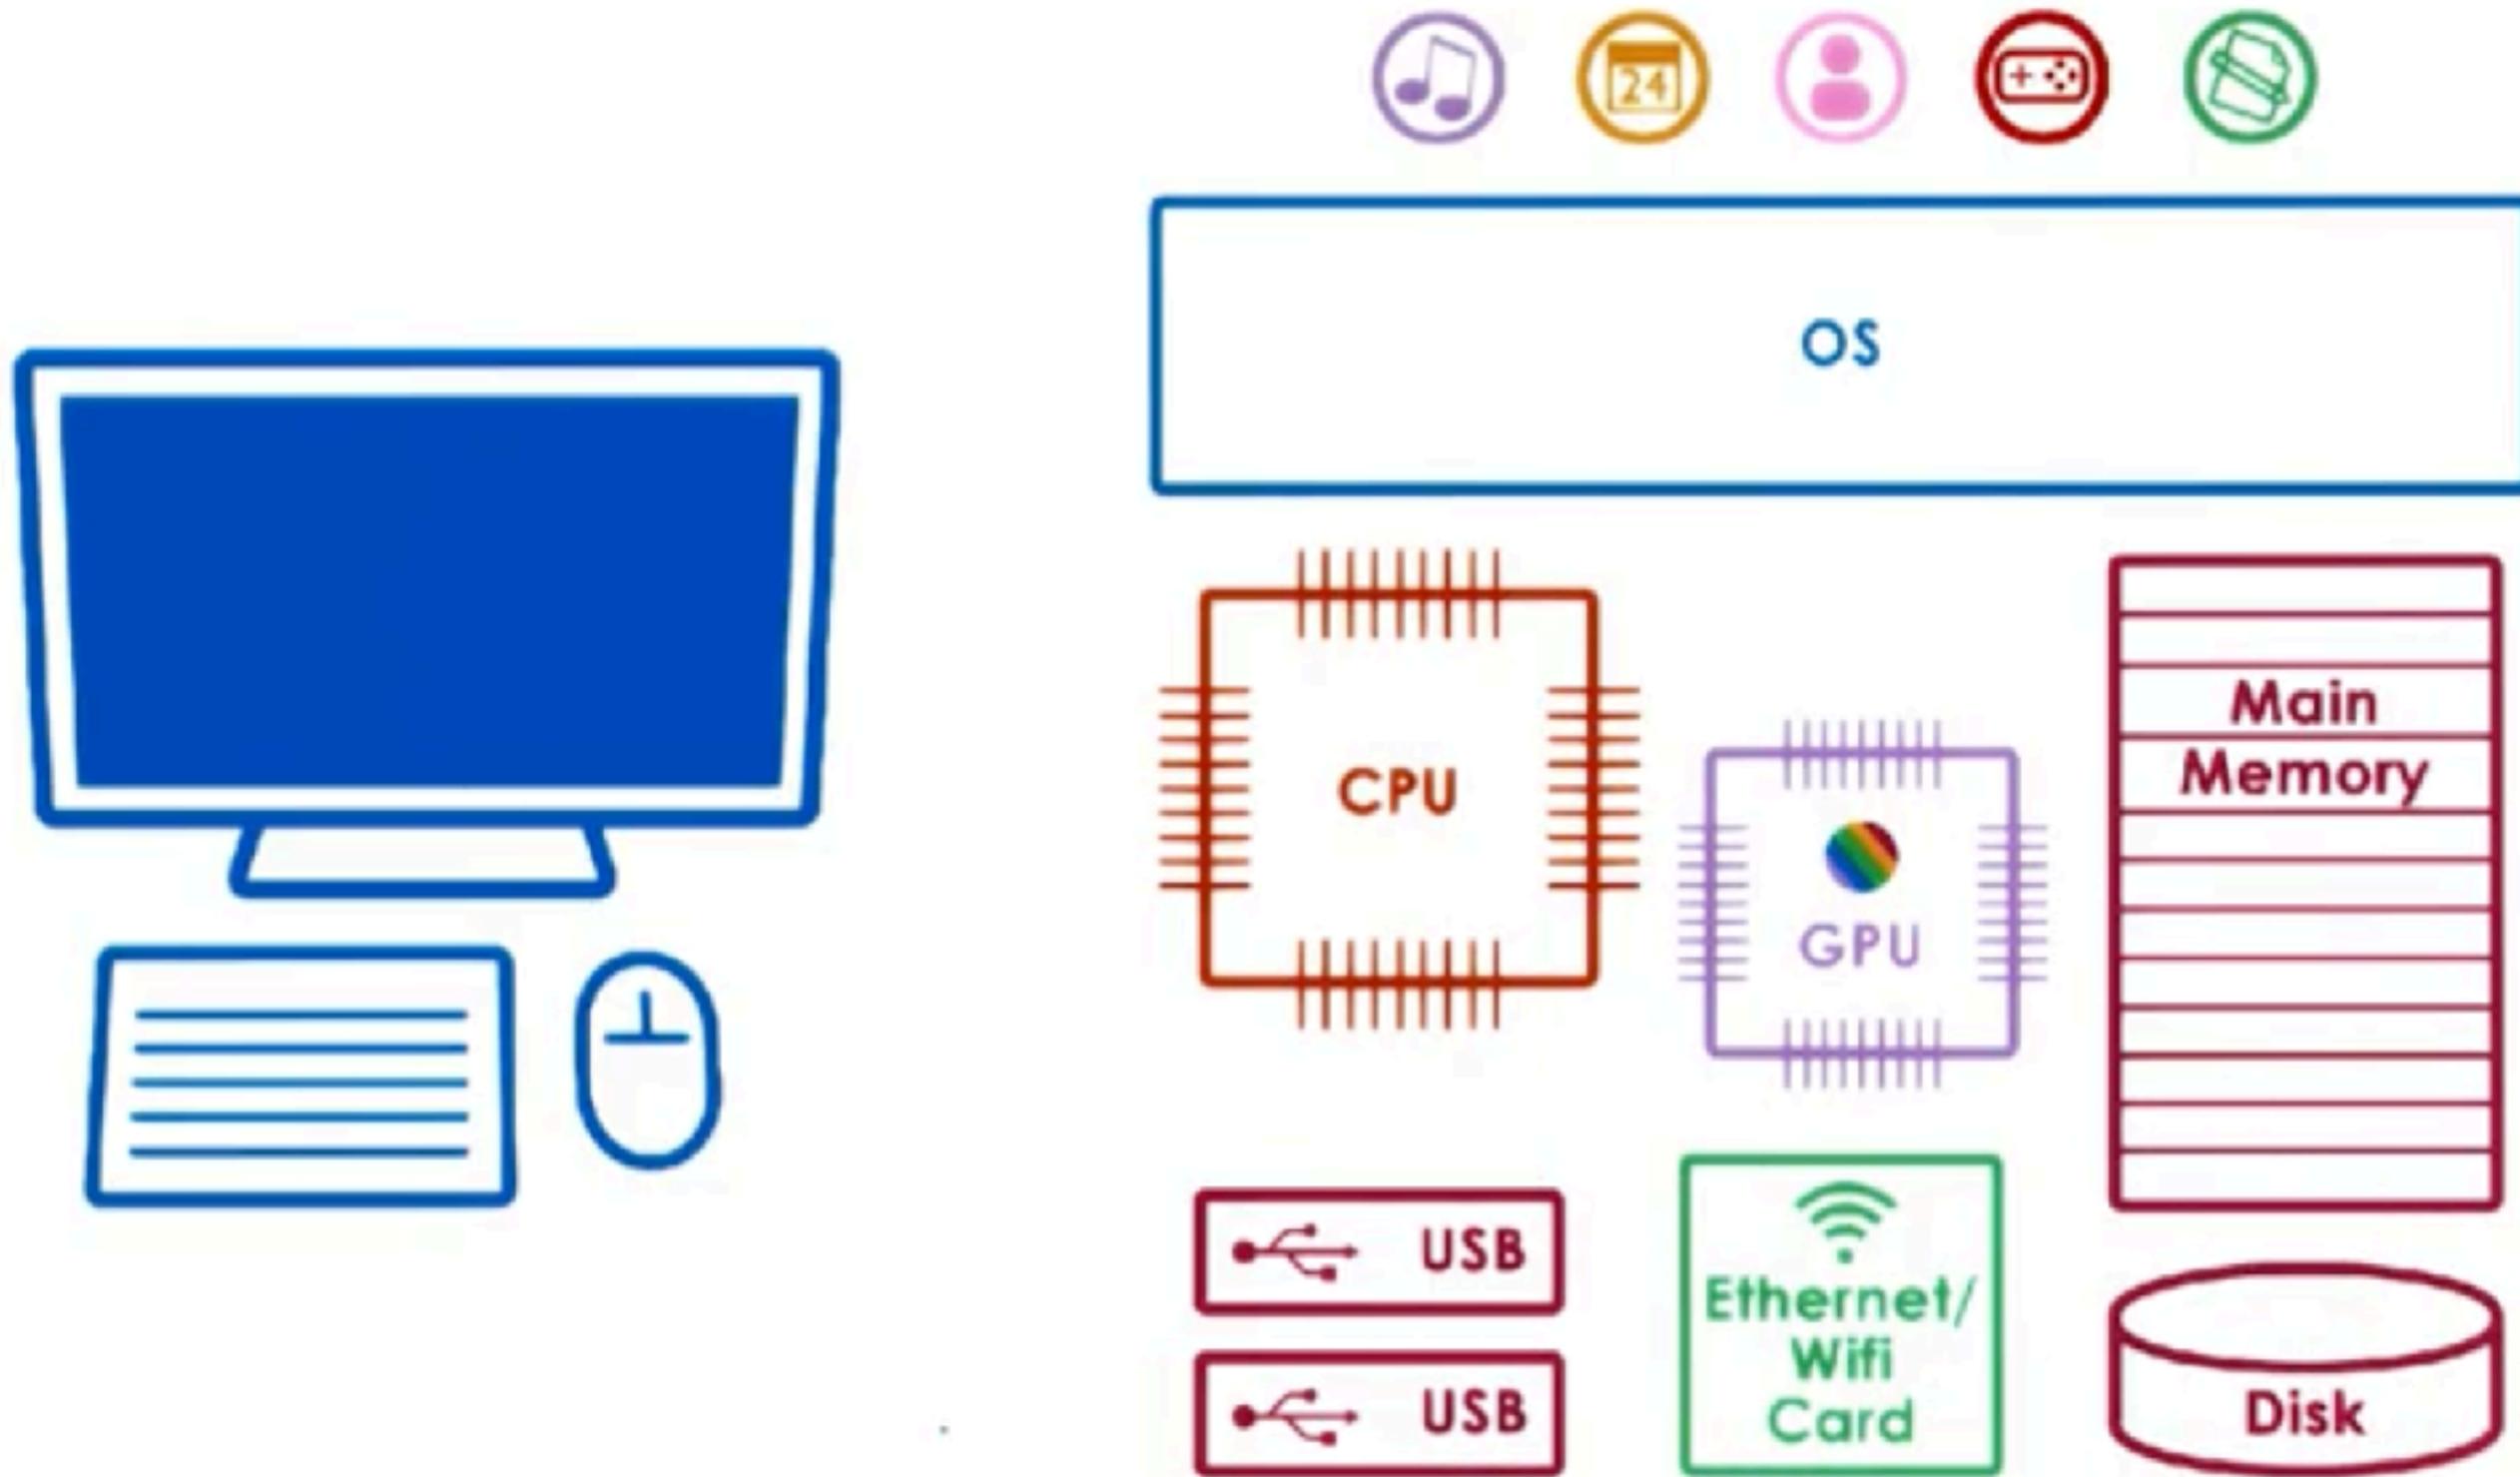


I/O Management

I/O Devices



basic I/O Device Features

Control registers

- command
- data transfers
- status

microcontroller == device's CPU

On device memory

Other logic

- e.g., analog to digital converters

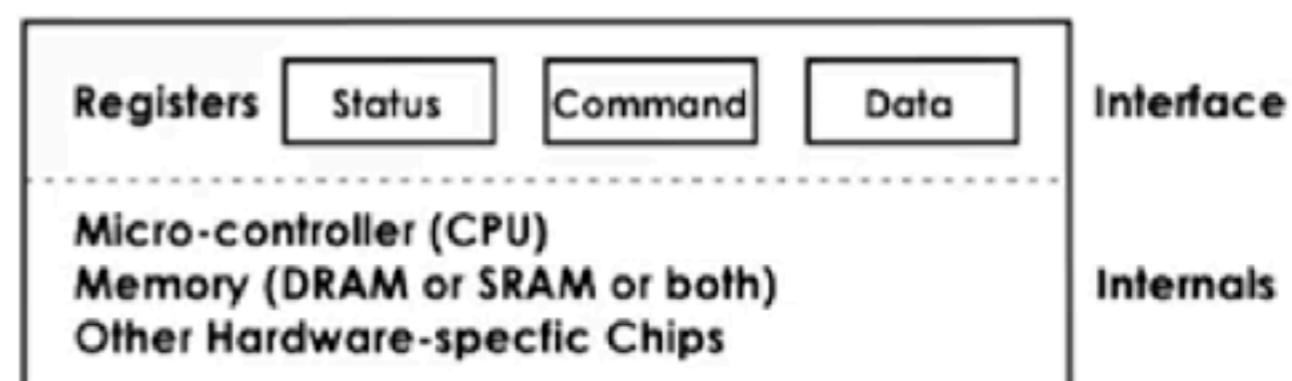
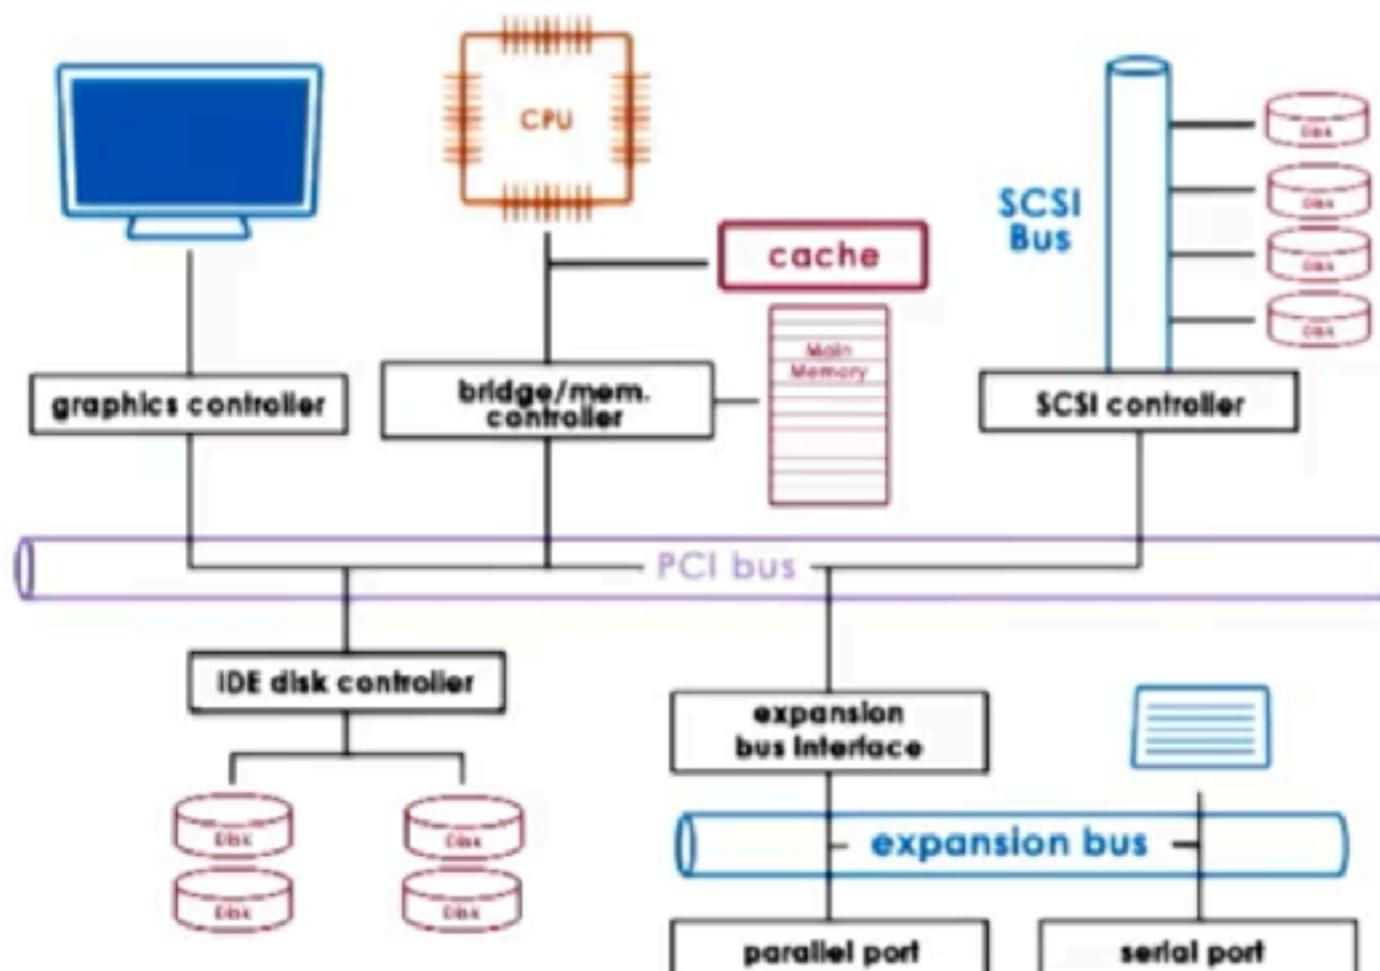


Figure 36.2: A Canonical Device

CPU - Device Interconnect



Peripheral Component
Interconnect (PCI)

- PCI Express (PCIe)
(> PCI-X > PCI)

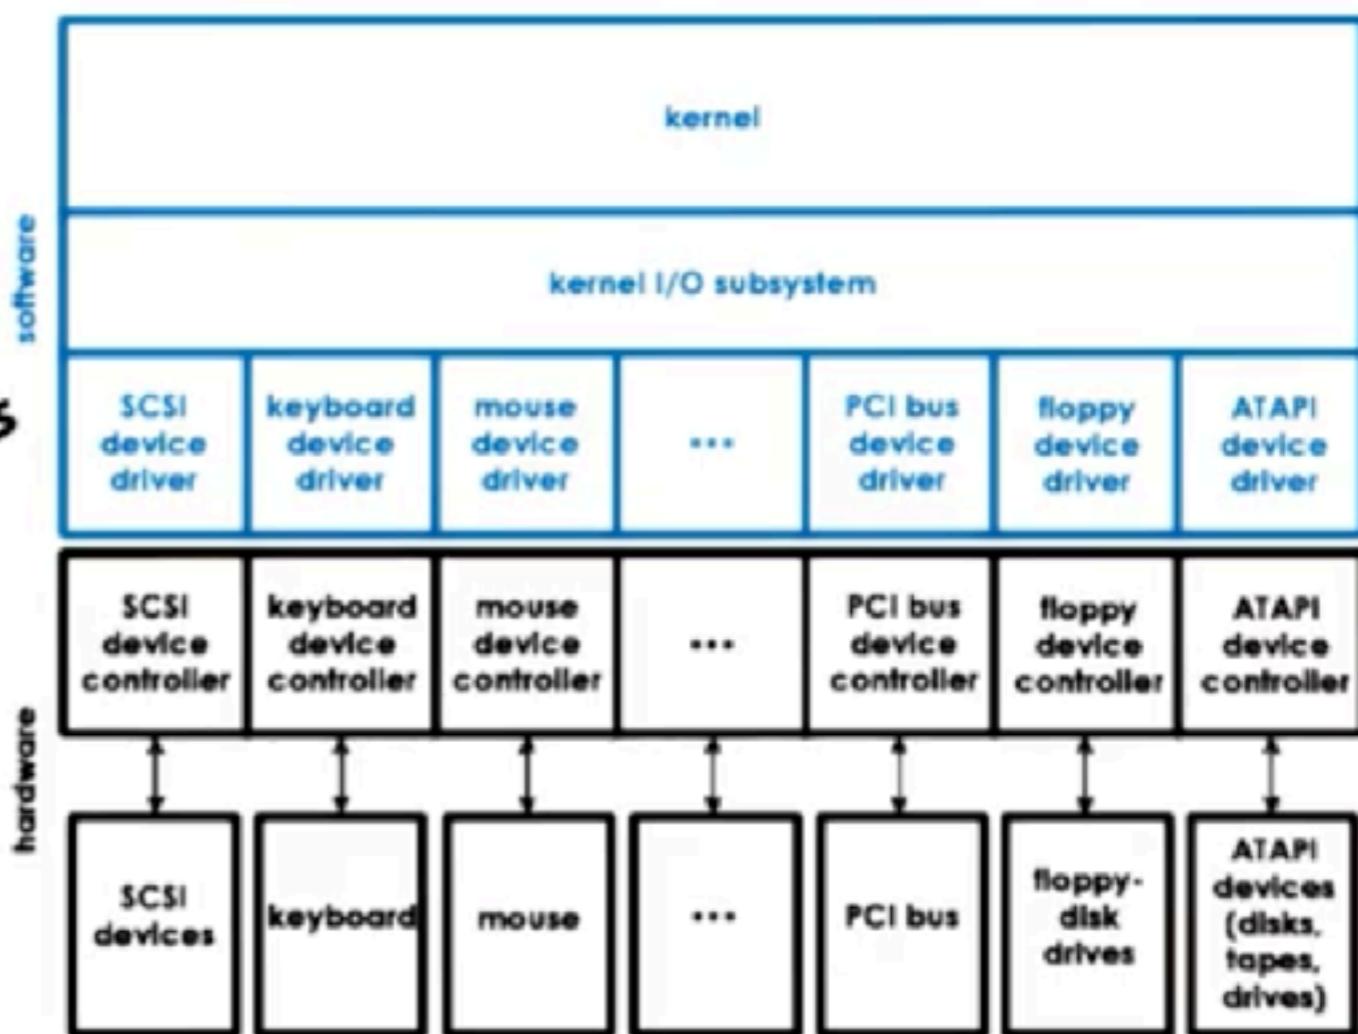
Other types of interconnects

- SCSI bus
- peripheral bus
- bridges handle differences

Device Drivers

Device Drivers

- per each device type
- responsible for device access, management and control
- provided by device manufacturers per OS/version
- each OS standardizes interfaces
 - device independence
 - device diversity



Device Types

Block: disk

- read / write blocks of data
- direct access to arbitrary block

OS representation of a device == special device file

Character : keyboard

- get / put character

Network devices

UNIX-like systems



- /dev
- tmpfs
- devfs



Pseudo Devices Quiz

Linux supports a number of pseudo ("virtual") devices that provide special functionality to a system. Given the following functions name the pseudo device that provides that functionality.

- accept and discard all output
(produces no output)

/dev/null

- produces a variable-length string
of pseudo-random numbers

/dev/random

Note: Answer by using the full path (e.g., /dev/loop0)



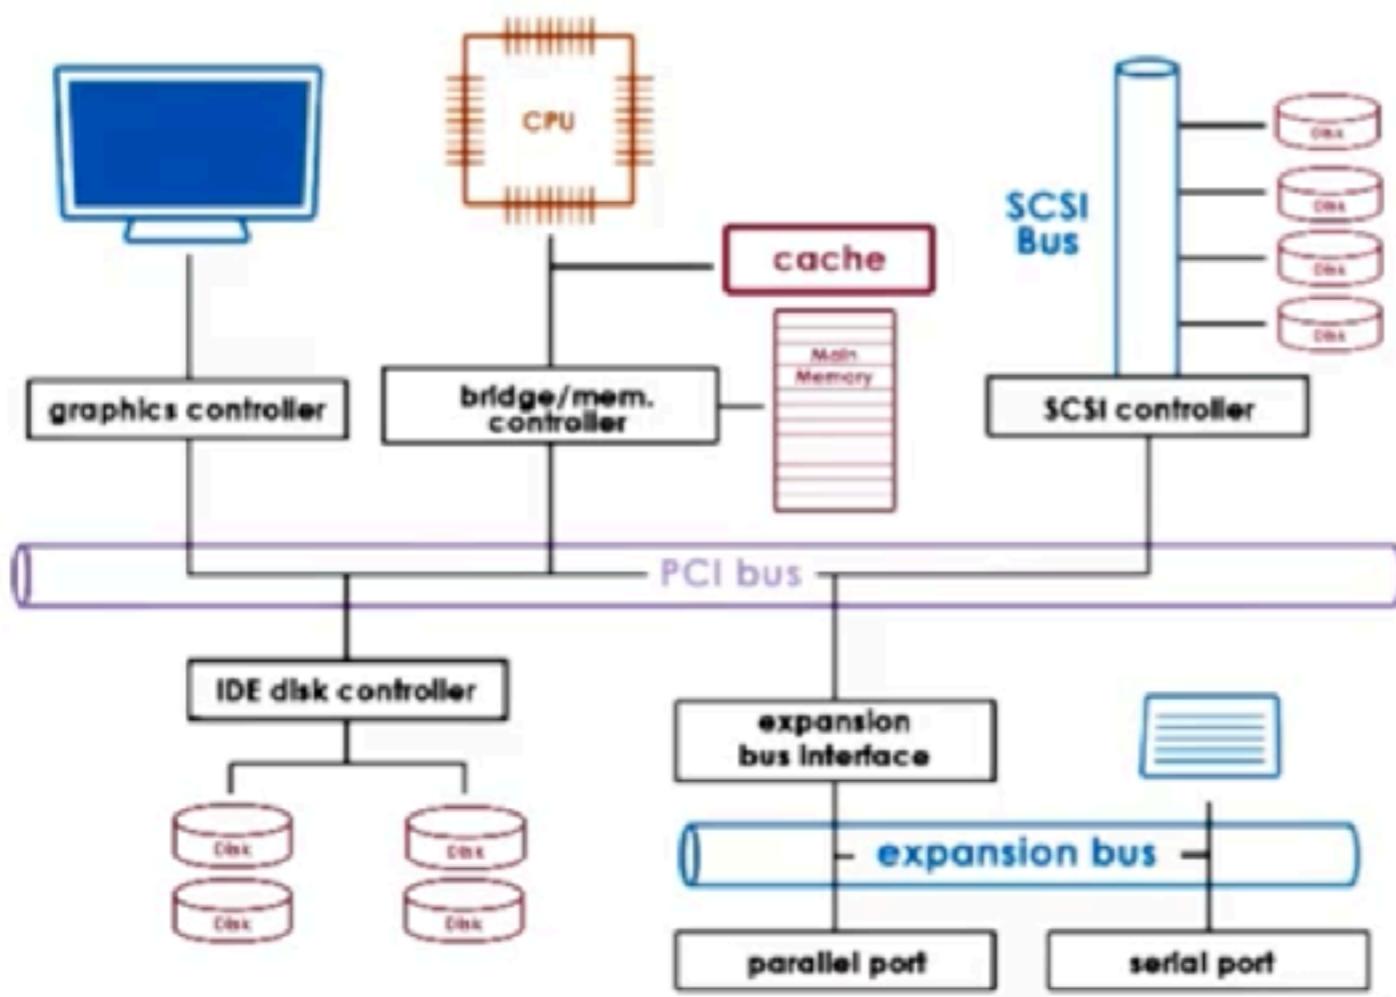
Looking at /Dev Quiz

Run the command 'ls -la /dev' in a Linux environment. What are some of the device names you see? Enter at least five device names in a comma-separated list in the text-box below.

hda, sda, tty, null, zero, ppp, lp,
mem, console, autoconf, ...

Note: You may use the Ubuntu VM provided for this course. See the Instructor Notes for a download link.

CPU - Device Interactions



access device registers
== memory load/store

memory-mapped I/O

- part of 'host' physical memory dedicated for device interactions
- Base Address Registers (BAR)

I/O port

- dedicated in/out instructions for device access
- target device (I/O port) and value in register

CPU/Device Interaction

Interrupt



interrupt handling
steps



can be generated as
soon as possible

Polling



when convenient for
OS



delay or
CPU overhead

Programmed I/O

no additional hardware support

CPU "programs" the device

- via command registers
- data movement

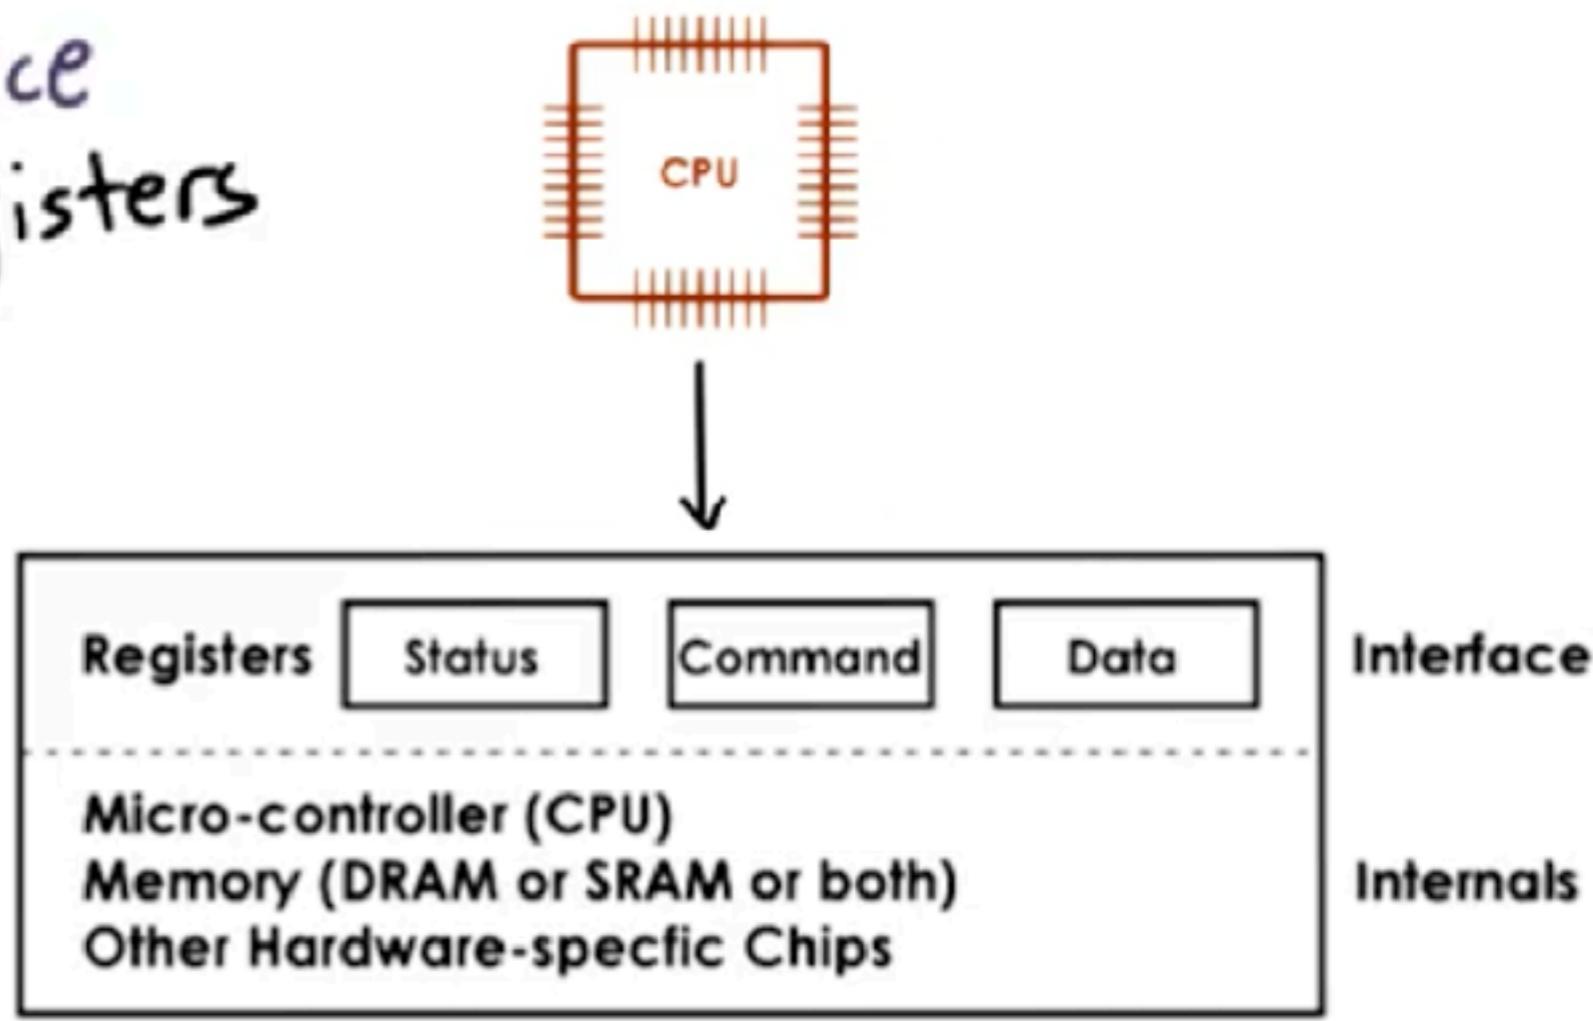


Figure 36.2: A Canonical Device

Programmed I/O

Example: NIC, data == network packet

- write command to request packet transmission
- copy packet to data registers
- repeat until packet sent

e.g., 1500 B packet; 8 byte regs or bus

=> 1 (for bus command) +
188 (for data)

=> 189 CPU store instructions

Direct Memory Access (DMA)

relies on DMA controller

CPU "programs" the device
- via command registers
- via DMA controls

For DMAs:

- data buffer must be in physical memory until transfer completes

=> pinning regions
(non-swappable)

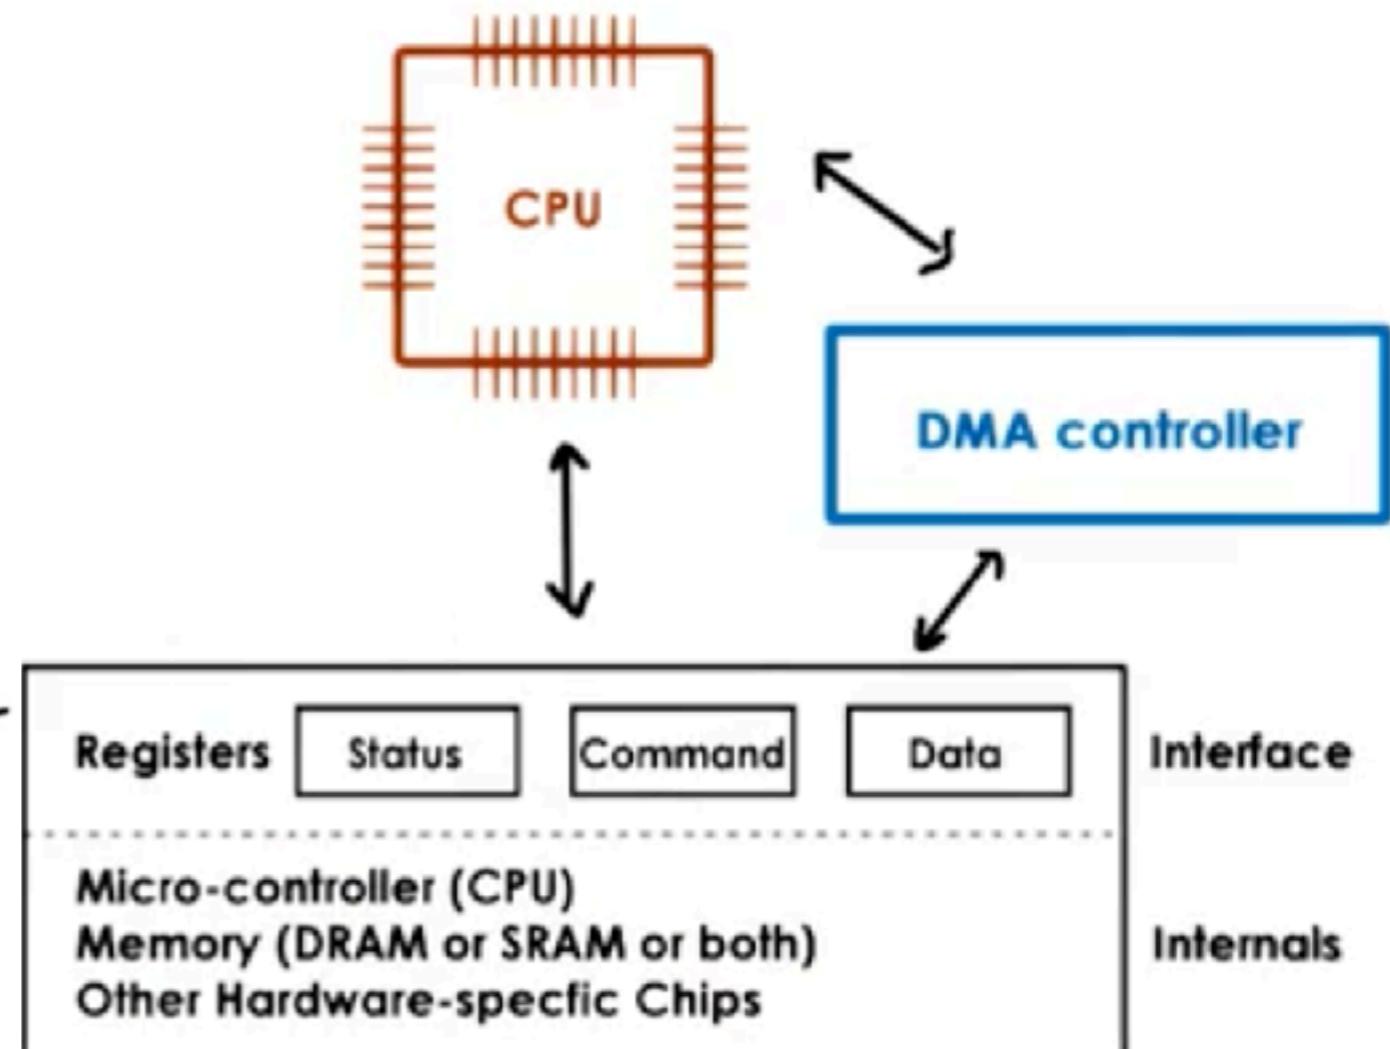


Figure 36.2: A Canonical Device

Example : NIC, data == network packet

- write command to request packet transmission
- configure DMA controller with in-memory address and size of packet buffer

e.g., 1500B w/ 8 byte regs or bus

=> | store instruction +
| DMA configure



DMA vs. PIO Quiz

For a hypothetical system, assume the following:

- it costs 1 cycle to run a store instruction to a device register
 - it costs 5 cycles to configure a DMA controller
 - the PCI-bus is 8 bytes wide
 - all devices support both DMA and PIO access
- Which device access method is best for the following devices?

Keyboard :

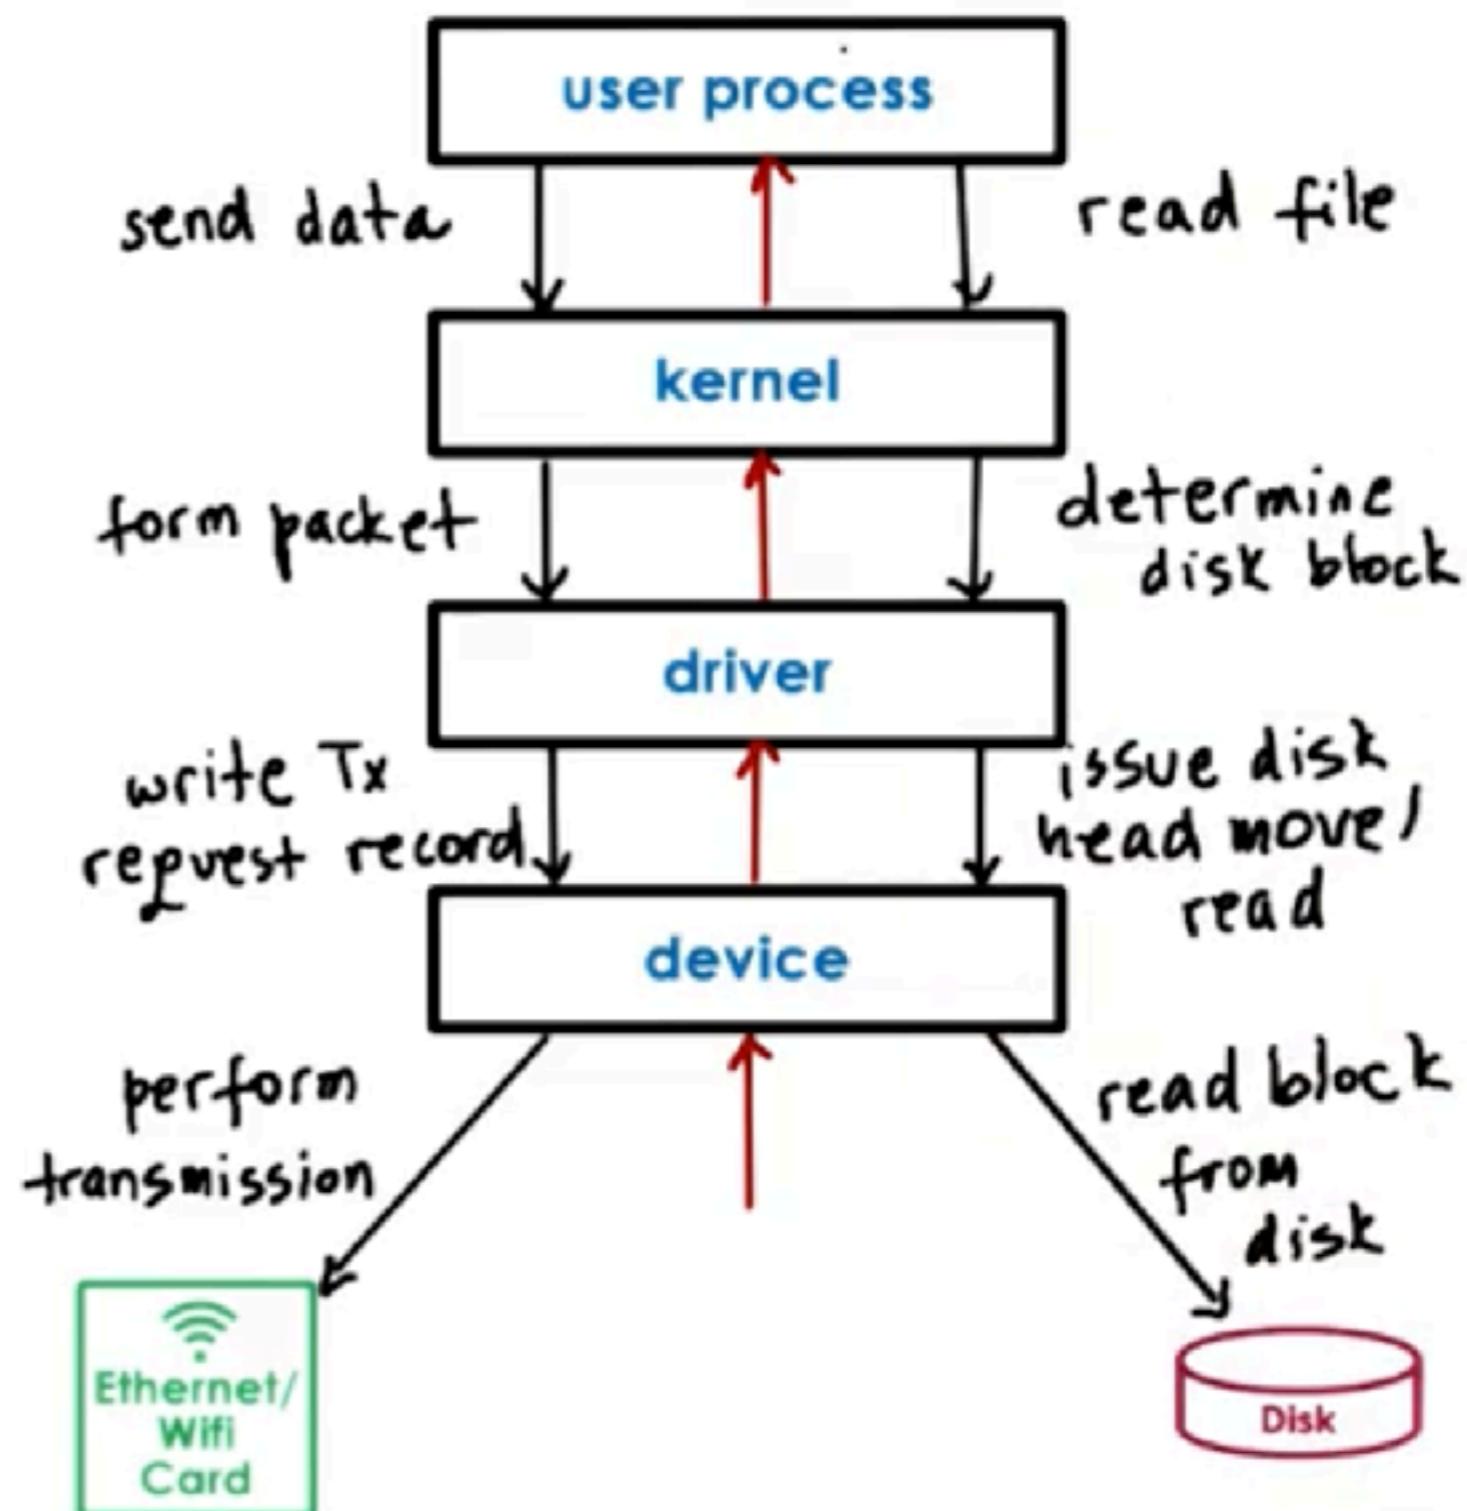
<input checked="" type="checkbox"/> PIO	<input type="checkbox"/> DMA	<input type="checkbox"/> Depends
---	------------------------------	----------------------------------

NIC :

<input type="checkbox"/> PIO	<input checked="" type="checkbox"/> DMA	<input checked="" type="checkbox"/> Depends
------------------------------	---	---

Typical Device Access

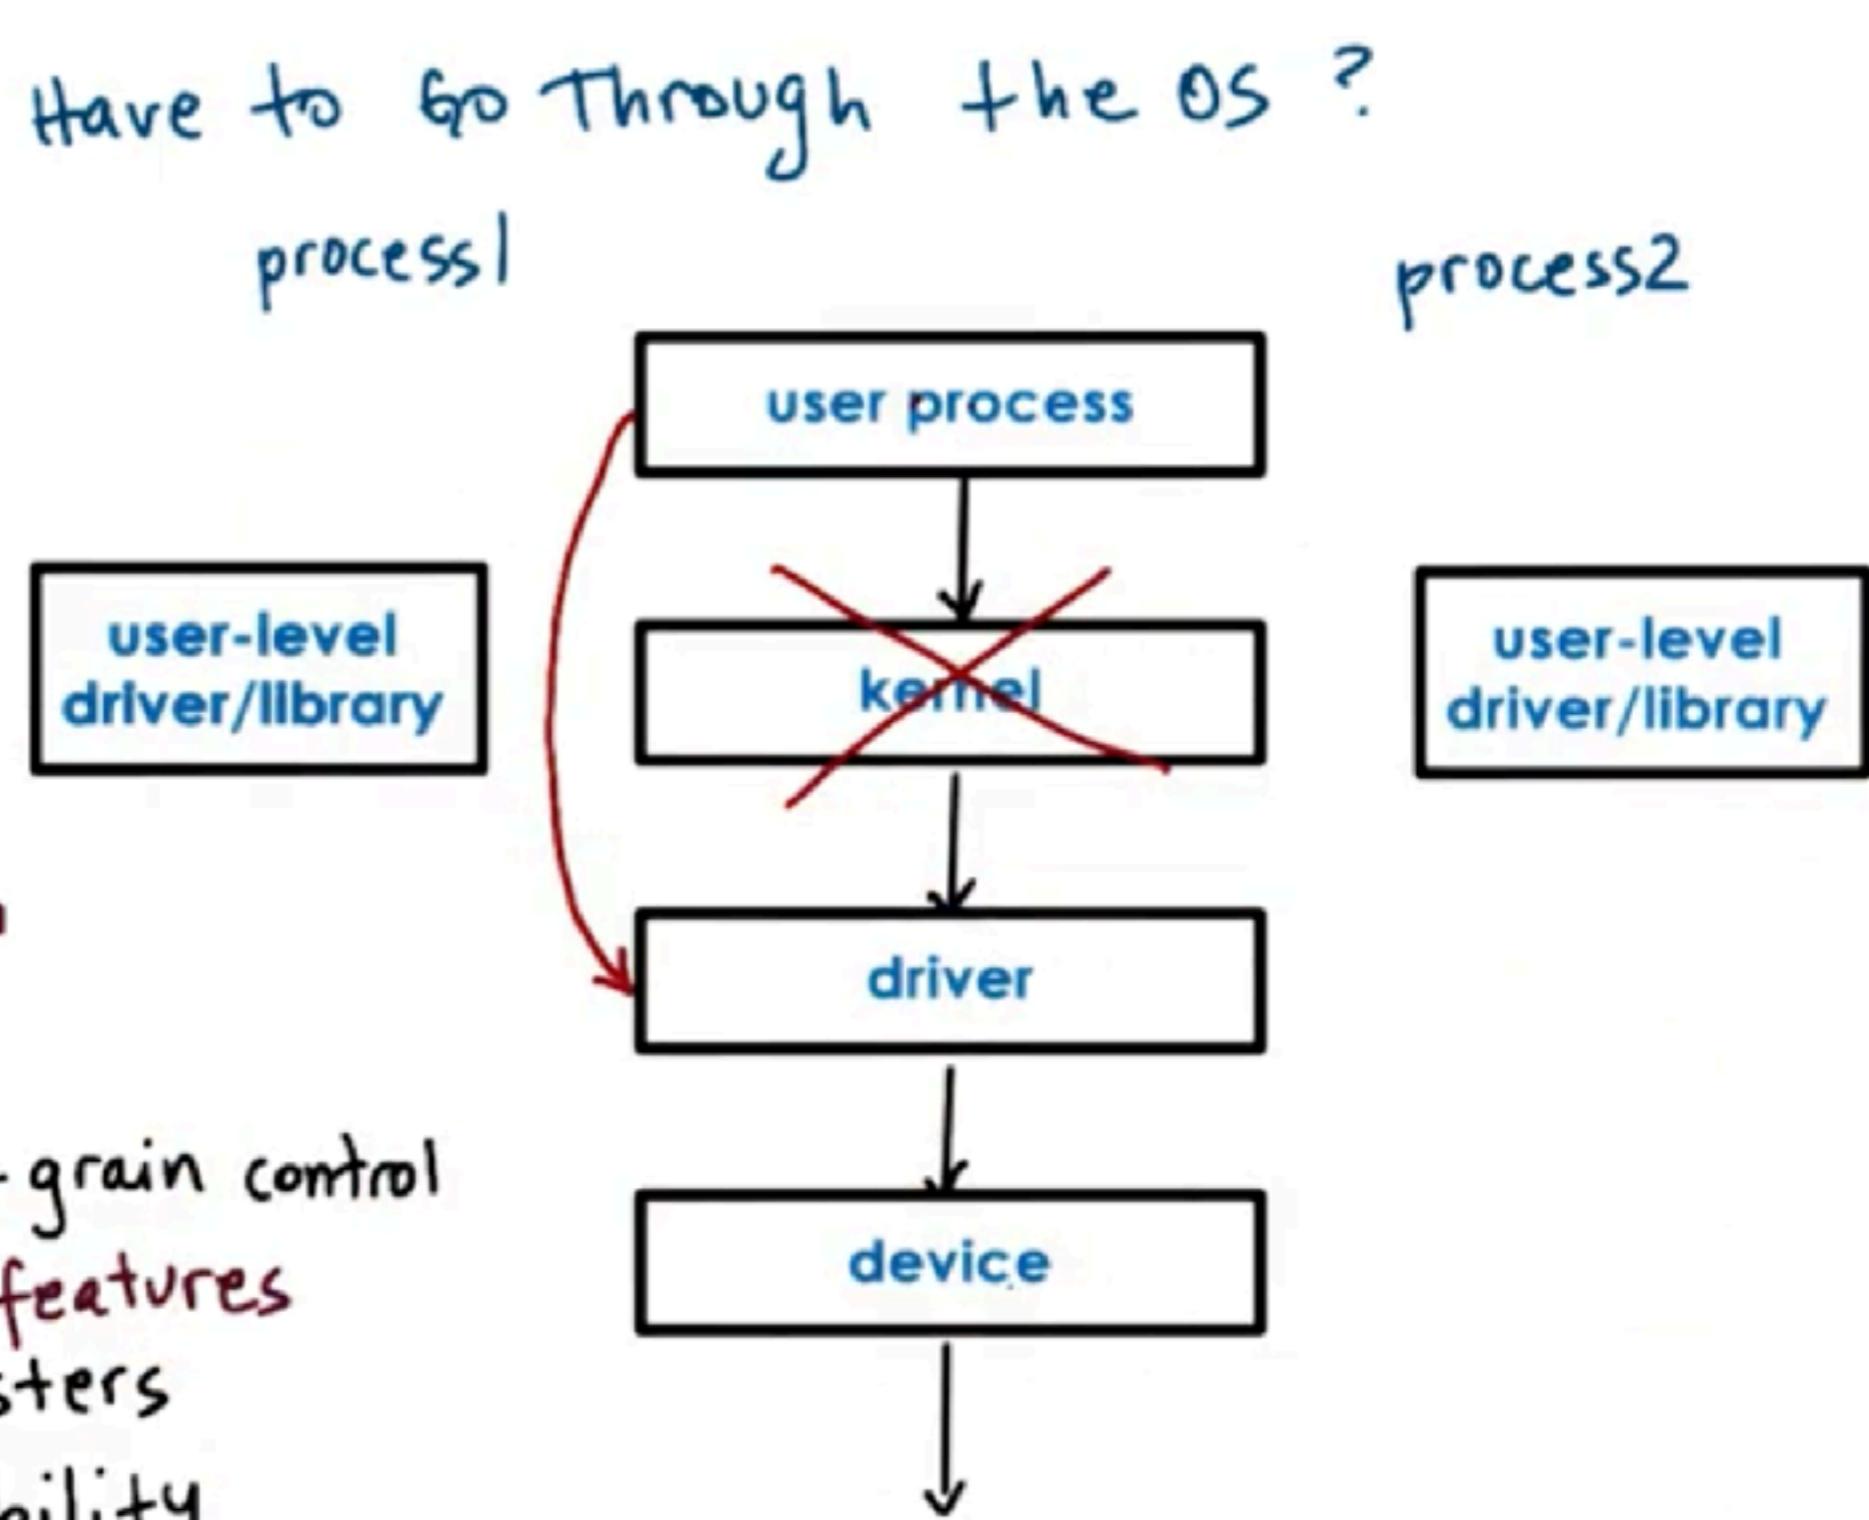
- system call
- in-kernel stack
- driver invocation
- device request configuration
- device performs request



Do You Have to Go Through the OS ?

OS Bypass

- device regs / data directly accessible
- OS configures then out-of-the way
- "user-level driver" (~ library)
- OS retains coarse-grain control
- relies on device features
 - sufficient registers
 - demux capability



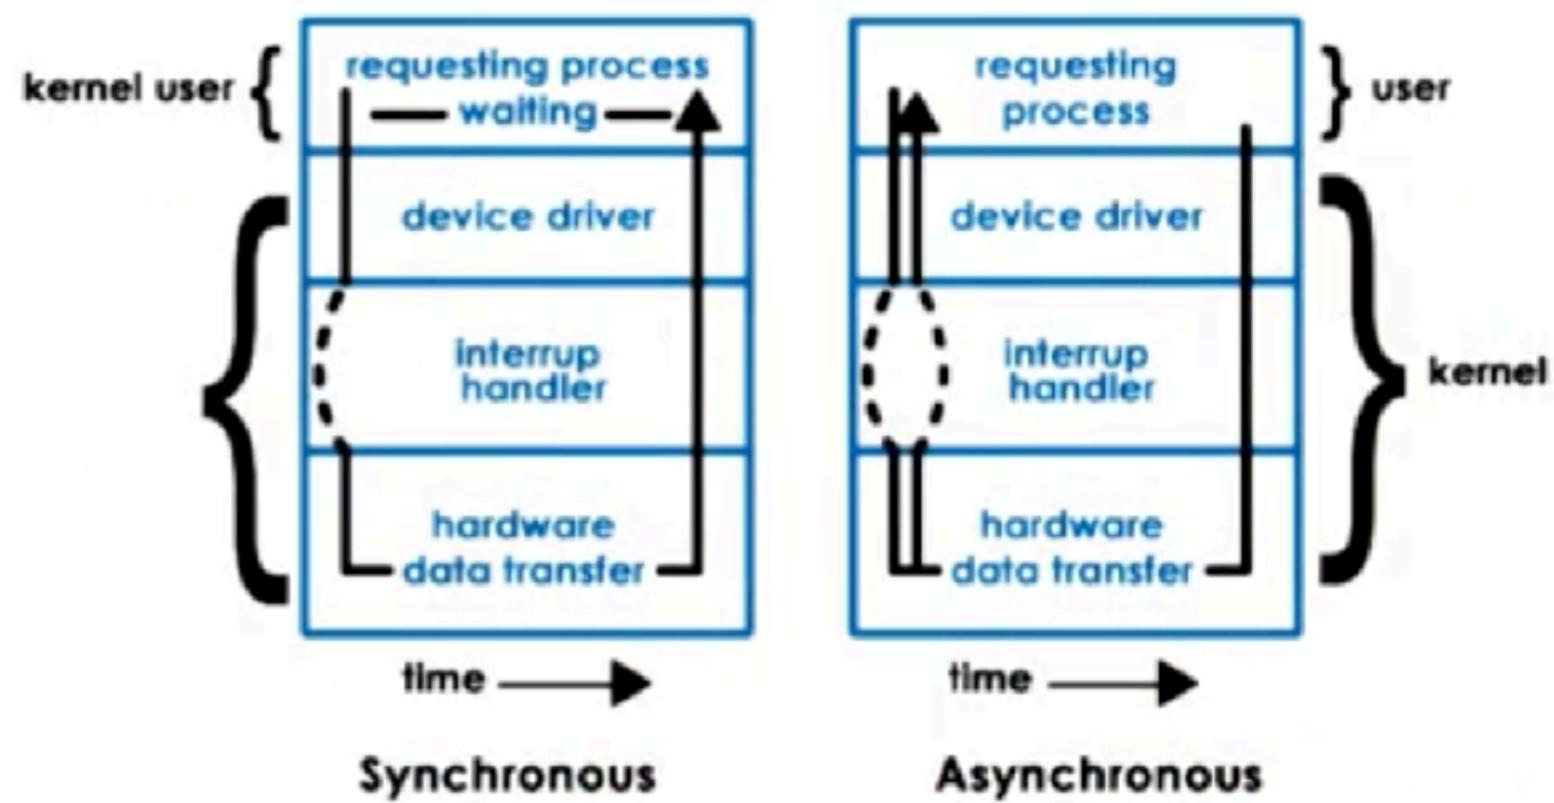
Sync vs Async I/O

Synchronous I/O operations
=> process blocks

Asynchronous I/O operations
=> process continues

Later ...

- process checks and retrieves result
- OR
- process is notified that the operation completed and results are ready



Block Device Stack

processes use files => logical storage unit

kernel file system (FS)

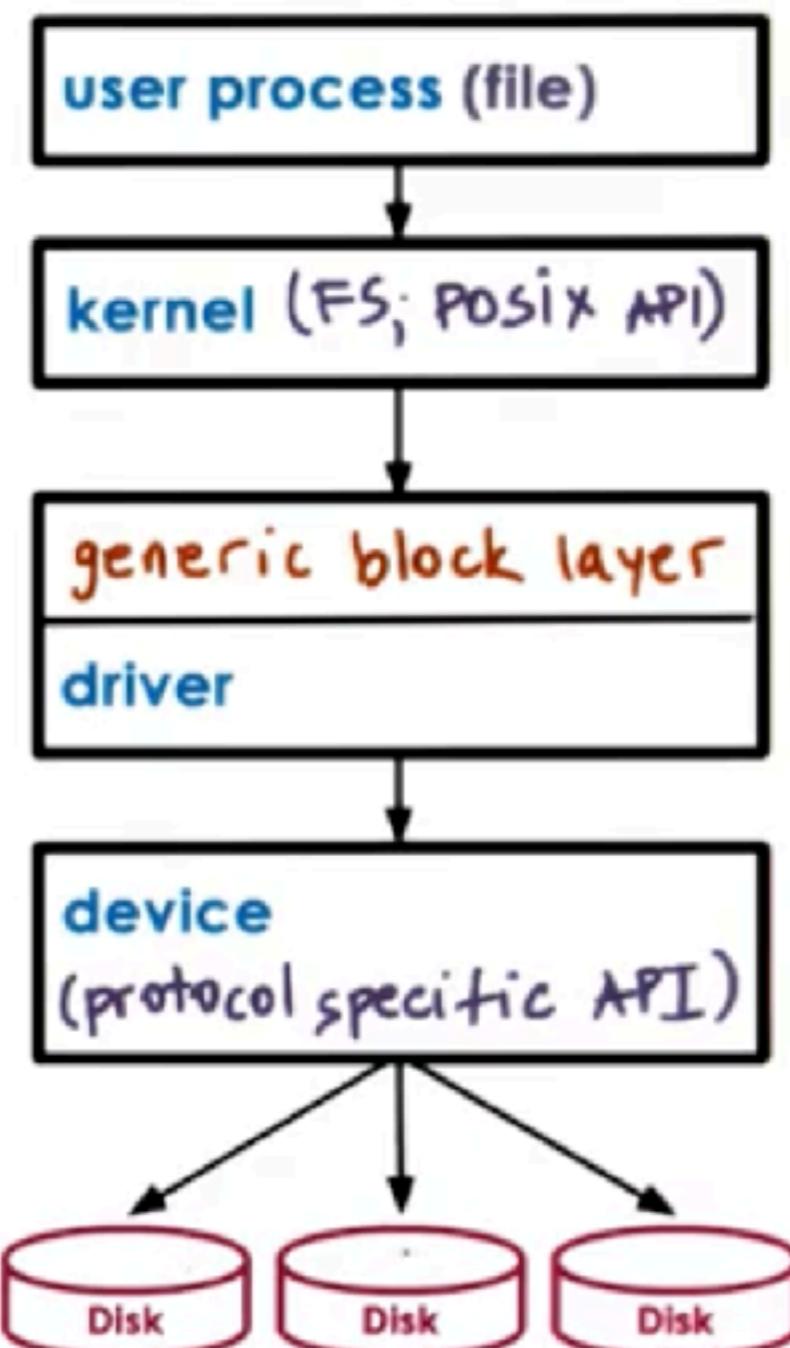
- where, how to find and access file
- OS specifies interface

generic block layer

- OS standardized block interface

device driver

block device
typical storage
for files





Block Device Quiz

In Linux, the `ioctl()` command can be used to manipulate devices.

Complete the code snippet, using `ioctl()`, to determine the size of a block device.

```
// ...
int fd;
unsigned long numblocks = 0;

fd = open(argv[1], O_RDONLY);
ioctl(fd, BLKGETSIZE, &numblocks);
close(fd);

// ...
```